

Problem Solving Workshop

Phase 1: Foundations

Instructor: Abhishek Bansal
Academic Year: 2025

Course Overview

This two-week workshop introduces foundational problem-solving techniques through classic puzzles, algorithmic strategies, and real-world modeling exercises. Students will build confidence in recursion, search algorithms, data structures, and collaborative design.

Phase 1 Curriculum Overview

1. Sudoku – Baseline Quiz

- Sudoku, logic-grid, and handshake-count puzzles ($N(N - 1)/2$)
- Survey: "How did you solve each?"

Outcome: Gauge current skills and problem-solving intuition.

2. N-Queens Backtracking

- Explain rules
- Code stub: `def place_queens(row, board): ...`
- Test on $N = 4$ and $N = 8$

Outcome: Hands-on recursion and pruning.

3. Sudoku Solver

- Constraint propagation and backtracking
- Fill in logic for provided Python function signature

Outcome: Reinforce state-space search.

4. Graphs via Word Ladders – Mazes

- Introduction to graphs (nodes, edges)
- BFS/DFS on word ladders
- Maze-path visualization

Outcome: Bridge puzzles to graph algorithms.

5. Group Challenge – Extensions

- Teams choose from N-Queens, Sudoku, word ladders, mazes, or optional puzzles (Lights Out, Josephus, Tower of Hanoi, Coin Change, Word Search)
- Sketch and pitch an improved solution

Outcome: Collaboration and peer teaching.

6. Queues – Real-World Simulation

- Live-code a bank-line simulator using a queue
- Stack demo: undo/redo

Outcome: Apply data structures in real-world contexts.

7. Game AI – Tic-Tac-Toe

- Provide stub: `def minimax(board, player): ...`
- Implement and play versus AI

Outcome: Expose decision-tree logic.

8. Event-Driven Modeling: Traffic / Queue

- Model traffic lights or a bank queue with events
- Discuss real-world modeling challenges

Outcome: Experience event-driven thinking.

9. Ideation Workshop

- Brainstorm campus/community problems (library booking, canteen queues)
- Structured voting on top ideas

Outcome: Nurture interdisciplinary vision.

10. Mini-Pitch – Reflection

- Three-minute team pitches with solution outlines
- Reflection: "What did we learn?"

Outcome: Build confidence and prepare for Phase 2 capstone.

Setup Instructions

Students should follow these steps to access materials:

- Sign up or log into Replit: replit.com
- Take the Quiz on Google Forms: [Click here](#)
- View the Progress Tracker: [Click here](#)

Puzzles & Baseline Quiz

Ice-Breaker Puzzles

Complete the following exercises in 20 minutes:

- a. **Sudoku (4×4 grid):** Fill each row, column, and 2×2 block with numbers 1–4 exactly once.
- b. **Logic Grid Puzzle:** Three students (A, B, C) each solved *one* of three puzzles (maze, riddle, sudoku-mini) in *one* of three times (5, 10, 15 minutes), and no two students chose the same puzzle or time.
 - A did **not** take 5 minutes.
 - The riddle was solved in 10 minutes.
 - C solved the maze.
 - The sudoku-mini was completed in the shortest time (5 minutes).

Determine who solved which puzzle and in what time.

- c. **Handshake Count:** In a group of 7, every pair shakes hands exactly once. How many total handshakes occur?

Baseline Quiz

Answer the following in the next 10 minutes:

1. Compute the handshake count for $N = 7$ and verify using the formula $\frac{N(N-1)}{2}$.
2. Write pseudocode for computing handshake count for a general N :

```
function handshakeCount(N):  
    # Each of N people shakes hands with N-1 others  
    # Divide by 2 to avoid double-counting  
    return N * (N - 1) / 2
```

3. Reflect: Which problem-solving heuristic did you use most (decomposition, pattern recognition, abstraction)?

Backtracking I – N-Queens

Problem Description

Place one queen in each row of an $N \times N$ chessboard so that no two queens attack each other. A queen attacks along its row, column, and both diagonals.

Recursive Pseudocode

Use the following stub to implement your solution:

```
function placeQueens(row, board):
    if row > N:
        printSolution(board)
        return
    for col = 1 to N:
        if isSafe(row, col, board):
            board[row] = col
            placeQueens(row + 1, board)
            board[row] = 0 # backtrack

function isSafe(r, c, board):
    for prevRow = 1 to r - 1:
        prevCol = board[prevRow]
        if prevCol == c or
            abs(prevCol - c) == abs(prevRow - r):
            return false
    return true
```

Try it online: Run the N-Queens demo on Replit

Exercise

- Implement the pseudocode and print one valid arrangement for $N = 4$.
- Extend your code to count all solutions for $N = 4$ and report the total.
- *(Optional)* Test your solver for $N = 8$ and note the number of solutions.

Outcome: Reinforce backtracking and recursion.

Sudoku Solver Session

Problem Description

Students will implement a backtracking algorithm to solve a 4×4 Sudoku. Each row, column, and 2×2 block must contain numbers 1–4 exactly once.

Recursive Backtracking Stub

Use this function signature:

```
function solveSudoku(grid):
    if no empty cells:
        printSolution(grid)
        return true
    pick an empty cell (r,c)
    for num = 1 to 4:
        if isValid(grid, r, c, num):
            grid[r][c] = num
            if solveSudoku(grid): return true
            grid[r][c] = 0    # backtrack
    return false
```

Try it online: Run the Sudoku Solver on Replit

Exercises

1. a. Write a helper function `findEmptyCell(grid)` that returns the coordinates (r, c) of the next empty cell (`grid[r][c] == 0`), or `null` if none remain.
 b. Fill in the body of `solveSudoku(grid)` using your `findEmptyCell` helper.
2. Run your solver on the sample 4×4 puzzle:

```
[ [0,2,0,4],
  [3,0,1,0],
  [0,1,0,3],
  [4,3,0,0] ]
```

3. **(Optional)** Extend your solver to handle full 9×9 puzzles.

Outcome: Reinforce backtracking and constraint propagation.

Graphs via Word Ladders & Mazes

Problem Description

Students will explore graphs (vertices and edges) and then apply graph-search to two puzzles:

- **Word Ladder**

Transform one word into another by changing exactly one letter at a time; each intermediate word must appear in the dictionary.

- **Maze Navigation**

Navigate a 2D ASCII-maze from a start (S) to an exit (E), moving only through open cells “.”.

Breadth-First Search (BFS) Stub

Use this pseudocode to find the shortest path in an unweighted graph (queue operations implicit):

```
function BFS(start, goal):
    Q ← empty queue
    enqueue(Q, start); mark start visited
    while Q not empty:
        v ← dequeue(Q)
        if v == goal: return path-to(v)
        for each neighbor u of v:
            if not visited(u):
                mark u visited; enqueue(Q, u)
    return "no path"
```

Depth-First Search (DFS) Stub

Use this pseudocode to traverse all reachable nodes:

```
function DFS(v):
    mark v visited
    for each neighbor u of v:
        if not visited(u):
            DFS(u)
```

Sample Inputs

- **Word Ladder**

```
beginWord = "hit"
endWord   = "cog"
dict      = {"hot", "dot", "dog", "lot", "log", "cog"}
```

- **Maze (ASCII)**

```
#####
#S..#.#.#
#.#.#.#.#
#..E##.##
#####
```

(Here S = start, E = exit, . = open, # = wall.)

Try it online: Run the BFS Word-Ladder demo on Replit

Try it online: Run the BFS Maze Navigation demo on Replit

Try it online: Run the DFS Word-Ladder demo on Replit

Try it online: Run the DFS Maze Navigation demo on Replit

Exercises

- 1 Implement **Word Ladder**: use the BFS stub to compute the shortest transformation from *beginWord* to *endWord*, returning the list of intermediate words.
- 2 **Maze-path Visualization**: represent the maze as a graph (cells as vertices, edges between adjacent open cells). Use BFS to find a shortest path from S→E, and output the path either as a list of coordinates or directional steps.
- 3 (*Optional*) Compare DFS vs BFS on the same maze: implement a DFS-based solver and observe whether it always finds the shortest path.

Outcome: Bridge puzzles to graph algorithms, reinforcing that both recursion (DFS) and queue-based search (BFS) generalize across problem domains.

Group Challenge — Extensions

After three days of “learning by doing” (N-Queens, Sudoku, Word Ladders, Mazes), today we’ll survey several classic puzzles and their high-level algorithms. You won’t write full code now, but you should:

1. See how each problem is posed.
2. Follow the pseudocode sketch.
3. Think of one small way you might extend or improve it.

Puzzles & Pseudocode Stubs

- **Lights Out**

A grid of lights toggles itself and its orthogonal neighbors when you press a cell. Goal: turn all lights off.

```
function solveLightsOut(state):
    if all cells are OFF:
        return solution # base case
    pick a cell (r,c)
    toggle(state, r, c)           # flip this cell + neighbors
    if solveLightsOut(state):
        return solution
    toggle(state, r, c)           # backtrack
    return failure
```

- **Josephus**

n people stand in a circle; every kth person is eliminated until one remains.

```
function josephus(n, k):
    if n == 1:
        return 0 # survivor index in zero-based
    prev = josephus(n-1, k)
    return (prev + k) mod n # re-index into current circle
```

- **Tower of Hanoi**

Move n disks among three pegs, one disk at a time, never placing larger atop smaller.

```
function hanoi(n, src, aux, dest):
    if n == 1:
        move disk from src to dest
        return
    hanoi(n-1, src, dest, aux) # move top n-1 to auxiliary
    move disk from src to dest
    hanoi(n-1, aux, src, dest) # move n-1 from auxiliary to dest
```

- **Coin Change (Dynamic)**

Count ways to make amount X using unlimited coins of given denominations.


```

function countWays(coins, X):
    dp = array[0..X] with dp[0] = 1
    for each coin in coins:
        for amt = coin to X:
            dp[amt] += dp[amt - coin]
    return dp[X]

```

- **Word Search**

Given a grid of letters, determine if a target word exists by moving through adjacent cells (up/down/left/right) without reuse.

```

function exists(word, grid):
    for each cell (r,c):
        if dfs((r,c), 0):
            return TRUE
    return FALSE

function dfs(pos, i):
    if i == length(word):
        return TRUE # full word found
    if grid[pos] != word[i]:
        return FALSE
    mark pos visited
    for each neighbor n of pos:
        if not visited(n) and dfs(n, i+1):
            return TRUE
    unmark pos
    return FALSE

```

Exercises

1. **Trace Lights Out.** On this 3×3 sample, show the first two presses (cells) you'd choose to begin solving:

```

1 0 0
0 1 0
0 0 1

```

2. **Compute Josephus.** For $n=7$, $k=3$, who survives? Show your calls and return values from the recurrence.
3. **(Optional)** Pick one of today's puzzles and sketch a small extension. For instance:

- Lights Out on a torus
- Variable k in Josephus
- Tower of Hanoi variants
- Greedy heuristics for Coin Change
- Word Search backtracking

Be ready to describe your idea in two sentences.

Outcome: Familiarize everyone with a broader family of recursive, dynamic-programming, and graph-search problems—and give you a chance to look under the hood of each algorithm, even if you’re not writing full code today.

Queues – Real-World Simulation

Problem Description:

Model a bank teller’s line as a queue: customers arrive, wait in FIFO order, and are served one at a time. We’ll also demo a stack-based undo/redo of teller actions.

Live-Code: Bank-Line Simulator

Use a queue to enqueue arriving customers and dequeue for service.

Stack Demo: Undo/Redo

Record each served customer on an `undoStack`. Allow “undo” (move from `undoStack` → `redoStack`) and “redo” (move back).

Pseudocode Stub

```
function serveBank(arrivals):
    Q = empty queue
    for customer in arrivals:
        enqueue(Q, customer)

    while not empty(Q):
        c = dequeue(Q)
        serve(c)
        push(undoStack, c)

function undo():
    if not empty(undoStack):
        c = pop(undoStack)
        unserve(c)
        push(redoStack, c)

function redo():
    if not empty(redoStack):
        c = pop(redoStack)
        serve(c)
        push(undoStack, c)
```

Visual Walkthrough

A quick state-diagram table (front→back):

Step	Queue (front→back)	undoStack	redoStack
start	[A, B, C, D]	[]	[]
serve A	[B, C, D]	[A]	[]
serve B	[C, D]	[A, B]	[]
undo	[C, D]	[A]	[B]
redo	[C, D]	[A, B]	[]

Complexity Note

- All queue operations (‘enqueue’/‘dequeue’) and stack operations (‘push’/‘pop’) run in $O(1)$ time per action.

Sample Input / Scenario

- Arrival order: [Alice, Bob, Carol, Dave]
- Actions:
 - Serve Alice, Bob, Carol
 - Undo twice
 - Redo once
- Track queue contents and stack states after each step.

Exercises

1. Given arrivals [Alice, Bob, Carol], serve two, undo one, then list:
 - Remaining queue
 - Contents of `undoStack` and `redoStack`
2. Extend the stub to support “VIP” arrivals that jump to the front of the queue. Sketch your code change.

Outcome:

Apply queues and stacks in a realistic simulation—reinforcing FIFO service and undo/redo mechanics.

Game AI – Tic-Tac-Toe

Problem Description:

Build an AI opponent for Tic-Tac-Toe using a decision-tree search (*minimax*). The AI enumerates possible move sequences and chooses the optimal one.

Pseudocode Stub

```
function minimax(board, player):
    if terminal(board):
        return utility(board)          # win / loss / draw score

    if player == MAX:                  # MAX = 'O'
        bestVal = -infty
        for move in legalMoves(board):
            val = minimax(apply(board, move), MIN)
            bestVal = max(bestVal, val)
        return bestVal
    else:                              # MIN = 'X'
        bestVal = +infty
        for move in legalMoves(board):
            val = minimax(apply(board, move), MAX)
            bestVal = min(bestVal, val)
        return bestVal

function bestMove(board, player):
    best = null
    bestVal = (player==MAX ? -infty : +infty)
    for move in legalMoves(board):
        val = minimax(apply(board, move), opposite(player))
        if (player==MAX and val > bestVal) or
            (player==MIN and val < bestVal):
            bestVal = val
            best = move
    return best
```

Exercises

1. **Trace minimax.** On the board below, X (= MIN) has just played; O (= MAX) has *one* legal move left—the centre square. Show the call sequence (*minimax* → *utility*) and the utility value finally returned.

X	O	X
O	–	O
X	O	X

2. **Implement utility.** Sketch pseudocode for `utility(board)` so that it returns `+1` if O wins, `–1` if X wins, and `0` for a draw.

Outcome:

Expose the mechanics of recursive decision-tree search and utility evaluation in simple game AI.

Event-Driven Modeling: Traffic / Queue

Problem Description:

Model a traffic-light controller (or bank queue) as an *event-driven* system. Cars/lights fire *events*—ARRIVAL or DEPARTURE—and a central loop dispatches them in time order.

Live Code – Event-Loop Simulator

Demonstrate an **event priority queue**¹ that holds pending events while a dispatcher invokes handlers.

Pseudocode Stub

```
function runSimulation(initialEvents):
    Q = empty priority-queue          # ordered by (time, D<A)
                                     # no built-in priority flag yet
    waitingList = []                 # current cars in line
    queueLength = 0                  # global length variable
    nextID = 1 + max(id for (_,_,id) in initialEvents)

    for e in initialEvents:
        enqueue(Q, e)

    while not empty(Q):
        (t, type, cid) = dequeue(Q)  # earliest event
        handleEvent(type, t, cid)

function handleEvent(type, t, cid):
    global nextID, queueLength

    if type == ARRIVAL:
        waitingList.append(cid)
        queueLength = len(waitingList)

        if queueLength == 1:          # line was empty
            schedule(DEPARTURE, t + serviceTime(), cid)

        schedule(ARRIVAL, t + arrivalInterval(), nextID)
        nextID += 1

    elif type == DEPARTURE:
        waitingList.pop(0)             # front car leaves
        queueLength = len(waitingList)

        if queueLength > 0:           # next car departs later
            nextFront = waitingList[0]
            schedule(DEPARTURE, t + serviceTime(), nextFront)
```

¹Tie rule: Events are ordered by (time, type) with DEPARTURE < ARRIVAL, i.e., departure events come before arrival events in the priority queue and are thus considered higher-priority.

```

updateSystemState(t, type, cid, queueLength)

function schedule(type, tNew, cid):
    event = (tNew, type, cid)    # plain 3-field tuple, ordered by time and type
    enqueue(Q, event)           # O(log n) priority-queue insert

function arrivalInterval(): return ⟨user-supplied⟩    # e.g. 3 s
function serviceTime():    return ⟨user-supplied⟩    # e.g. 4 s

function updateSystemState(t, type, cid, queueLen):
    log(t, type, cid, queueLen)    # placeholder|stats / GUI / etc.

```

Sample Input / Scenario

- **Warm-up** events: $[(0, \text{ARRIVAL}, c_1), (0, \text{ARRIVAL}, c_2)]$
- **Exercise set** arrivals: $[0, 2, 5]$
- **Handlers**
 - **ARRIVAL** – enqueue next arrival, *increment* queueLength
 - **DEPARTURE** – enqueue next departure, *decrement* queueLength
- Track both the *event queue* Q and the *system state* ($\text{queueLength} = |\text{waitingList}|$) at every step.

Exercises

1. **Trace event loop.** Assume $\text{arrivalInterval} = 6\text{ s}$ and $\text{serviceTime} = 4\text{ s}$. For arrivals at $[0, 2, 5]$, list the first five rows of

$$(\text{step}, Q, \text{dequeued } (t, \text{event}), \text{waitingList}, \text{queueLength})$$
 and write a short note for each row.
2. **PRIORITY event.** Extend `schedule` to accept a priority flag (e.g. extra tuple field) so that a “PRIORITY” event jumps to the front.

Outcome:

Understand event-driven architecture—event queues, dispatch loops, dynamic scheduling, and real-time queue-length tracking.

Ideation Workshop

Problem Description

Working in small teams, you will *discover* and *define* a real bottleneck on campus or in the local community—e.g. library room-booking clashes, canteen queue congestion, or hostel laundry scheduling. The aim is to express the problem precisely and decide which issues merit an algorithmic solution.

Live Demonstration – Minimal Vote-Tally Script

The facilitator walks through a 10-line Python function that counts dot-votes and produces a ranked list of ideas.

```
# ideas    : list[str]
# ballots  : list[list[int]] # indices voted by each student
# Each student can cast up to 3 dots.

def tally_votes(ideas, ballots, dots_per_student=3):
    scores = {i: 0 for i in range(len(ideas))}
    for ballot in ballots:
        assert len(ballot) <= dots_per_student
        for idx in ballot:
            scores[idx] += 1      # one dot = one point
    return sorted(scores.items(), key=lambda p: -p[1])
```

Pseudocode Stub – Workshop Helper

```
function runIdeationWorkshop(studentIdeas):
    clusters = cluster_similar(studentIdeas)    # manual grouping / DBSCAN
    display_table(clusters)                     # show on projector

    votes    = collect_dot_votes(clusters)      # 3 dots per learner
    ranking  = sort_by_votes(votes)

    topIdea  = ranking[0]
    return draft_problem_statement(topIdea)
```

Warm-Up – Everyday Bottlenecks

Discuss in pairs: which campus processes could benefit from an algorithmic or data-driven approach?

1. Library *seat* / *room* booking clashes
2. Canteen queue congestion during peak breaks
3. Hostel laundry-machine scheduling fairness
4. Event ticket allocation transparency

Guided Brainstorm (15 min)

- **Individually:** Jot down as many campus/community problems as you can.
- **Teams of 3–4:** Cluster similar ideas and assign a concise label.

Structured Voting (5 min)

Each participant receives **three** dots/stickers. Distribute them among the clusters you believe are most pressing.

Idea / Cluster	Votes

Team Deliverable (10 min)

For the top-ranked idea, craft a concise problem statement covering:

- **Problem** – What exactly happens and why is it undesirable?
- **Stakeholders** – Who is affected and how?
- **Evidence** – A metric or anecdote illustrating scale or urgency.
- **Impact Goal** – Define success in measurable terms.

Outcome:

Broaden perspective beyond individual algorithms, nurturing interdisciplinary vision and opportunity discovery.