

Problem Solving Workshop

Phase 1: Foundations

Instructor: Abhishek Bansal
Academic Year: 2025

Course Overview

This two-week workshop introduces foundational problem-solving techniques through classic puzzles, algorithmic strategies, and real-world modeling exercises. Students will build confidence in recursion, search algorithms, data structures, and collaborative design.

Phase 1 Curriculum Overview

1. Sudoku – Baseline Quiz

- Sudoku, logic-grid, and handshake-count puzzles ($N(N - 1)/2$)
- Survey: "How did you solve each?"

Outcome: Gauge current skills and problem-solving intuition.

2. N-Queens Backtracking

- Explain rules
- Code stub: `def place_queens(row, board): ...`
- Test on $N = 4$ and $N = 8$

Outcome: Hands-on recursion and pruning.

3. Sudoku Solver

- Constraint propagation and backtracking
- Fill in logic for provided Python function signature

Outcome: Reinforce state-space search.

4. Graphs via Word Ladders – Mazes

- Introduction to graphs (nodes, edges)
- BFS/DFS on word ladders
- Maze-path visualization

Outcome: Bridge puzzles to graph algorithms.

5. Group Challenge – Extensions

- Teams choose from N-Queens, Sudoku, word ladders, mazes, or optional puzzles (Lights Out, Josephus, Tower of Hanoi, Coin Change, Word Search)
- Sketch and pitch an improved solution

Outcome: Collaboration and peer teaching.

6. Queues – Real-World Simulation

- Live-code a bank-line simulator using a queue
- Stack demo: undo/redo

Outcome: Apply data structures in real-world contexts.

7. Game AI – Tic-Tac-Toe

- Provide stub: `def minimax(board, player): ...`
- Implement and play versus AI

Outcome: Expose decision-tree logic.

8. Event-Driven Modeling: Traffic / Queue

- Model traffic lights or a bank queue with events
- Discuss real-world modeling challenges

Outcome: Experience event-driven thinking.

9. Ideation Workshop

- Brainstorm campus/community problems (library booking, canteen queues)
- Structured voting on top ideas

Outcome: Nurture interdisciplinary vision.

10. Mini-Pitch – Reflection

- Three-minute team pitches with solution outlines
- Reflection: "What did we learn?"

Outcome: Build confidence and prepare for Phase 2 capstone.

Setup Instructions

Students should follow these steps to access materials:

- Sign up or log into Replit: replit.com
- Take the Quiz on Google Forms: [Click here](#)
- View the Progress Tracker: [Click here](#)

Puzzles & Baseline Quiz

Ice-Breaker Puzzles

Complete the following exercises in 20 minutes:

- a. **Sudoku (4×4 grid):** Fill each row, column, and 2×2 block with numbers 1–4 exactly once.
- b. **Logic Grid Puzzle:** Three students (A, B, C) each solved *one* of three puzzles (maze, riddle, sudoku-mini) in *one* of three times (5, 10, 15 minutes), and no two students chose the same puzzle or time.
 - A did **not** take 5 minutes.
 - The riddle was solved in 10 minutes.
 - C solved the maze.
 - The sudoku-mini was completed in the shortest time (5 minutes).

Determine who solved which puzzle and in what time.

- c. **Handshake Count:** In a group of 7, every pair shakes hands exactly once. How many total handshakes occur?

Baseline Quiz

Answer the following in the next 10 minutes:

1. Compute the handshake count for $N = 7$ and verify using the formula $\frac{N(N-1)}{2}$.
2. Write pseudocode for computing handshake count for a general N :

```
function handshakeCount(N):  
    # Each of N people shakes hands with N-1 others  
    # Divide by 2 to avoid double-counting  
    return N * (N - 1) / 2
```

3. Reflect: Which problem-solving heuristic did you use most (decomposition, pattern recognition, abstraction)?

Backtracking I – N-Queens

Problem Description

Place one queen in each row of an $N \times N$ chessboard so that no two queens attack each other. A queen attacks along its row, column, and both diagonals.

Recursive Pseudocode

Use the following stub to implement your solution:

```
function placeQueens(row, board):
    if row > N:
        printSolution(board)
        return
    for col = 1 to N:
        if isSafe(row, col, board):
            board[row] = col
            placeQueens(row + 1, board)
            board[row] = 0 # backtrack

function isSafe(r, c, board):
    for prevRow = 1 to r - 1:
        prevCol = board[prevRow]
        if prevCol == c or
            abs(prevCol - c) == abs(prevRow - r):
            return false
    return true
```

Try it online: Run the N-Queens demo on Replit

Exercise

- Implement the pseudocode and print one valid arrangement for $N = 4$.
- Extend your code to count all solutions for $N = 4$ and report the total.
- *(Optional)* Test your solver for $N = 8$ and note the number of solutions.

Outcome: Reinforce backtracking and recursion.

Sudoku Solver Session

Problem Description

Students will implement a backtracking algorithm to solve a 4×4 Sudoku. Each row, column, and 2×2 block must contain numbers 1–4 exactly once.

Recursive Backtracking Stub

Use this function signature:

```
function solveSudoku(grid):
    if no empty cells:
        printSolution(grid)
        return true
    pick an empty cell (r,c)
    for num = 1 to 4:
        if isValid(grid, r, c, num):
            grid[r][c] = num
            if solveSudoku(grid): return true
            grid[r][c] = 0    # backtrack
    return false
```

Try it online: Run the Sudoku Solver on Replit

Exercises

1. a. Write a helper function `findEmptyCell(grid)` that returns the coordinates (r, c) of the next empty cell (`grid[r][c] == 0`), or `null` if none remain.
b. Fill in the body of `solveSudoku(grid)` using your `findEmptyCell` helper.
2. Run your solver on the sample 4×4 puzzle:

```
[ [0,2,0,4],
  [3,0,1,0],
  [0,1,0,3],
  [4,3,0,0] ]
```

3. **(Optional)** Extend your solver to handle full 9×9 puzzles.

Outcome: Reinforce backtracking and constraint propagation.

Graphs via Word Ladders & Mazes

Problem Description

Students will explore graphs (vertices and edges) and then apply graph-search to two puzzles:

- **Word Ladder**

Transform one word into another by changing exactly one letter at a time; each intermediate word must appear in the dictionary.

- **Maze Navigation**

Navigate a 2D ASCII-maze from a start (S) to an exit (E), moving only through open cells “.”.

Breadth-First Search (BFS) Stub

Use this pseudocode to find the shortest path in an unweighted graph (queue operations implicit):

```
function BFS(start, goal):
    Q ← empty queue
    enqueue(Q, start); mark start visited
    while Q not empty:
        v ← dequeue(Q)
        if v == goal: return path-to(v)
        for each neighbor u of v:
            if not visited(u):
                mark u visited; enqueue(Q, u)
    return "no path"
```

Depth-First Search (DFS) Stub

Use this pseudocode to traverse all reachable nodes:

```
function DFS(v):
    mark v visited
    for each neighbor u of v:
        if not visited(u):
            DFS(u)
```

Sample Inputs

- **Word Ladder**

```
beginWord = "hit"
endWord   = "cog"
dict      = {"hot", "dot", "dog", "lot", "log", "cog"}
```

- **Maze (ASCII)**

```
#####
#S..#.#.#
#.#.#.#.#
#..E##.##
#####
```

(Here S = start, E = exit, . = open, # = wall.)

Try it online: Run the BFS Word-Ladder demo on Replit

Try it online: Run the BFS Maze Navigation demo on Replit

Try it online: Run the DFS Word-Ladder demo on Replit

Try it online: Run the DFS Maze Navigation demo on Replit

Exercises

- 1 Implement **Word Ladder**: use the BFS stub to compute the shortest transformation from *beginWord* to *endWord*, returning the list of intermediate words.
- 2 **Maze-path Visualization**: represent the maze as a graph (cells as vertices, edges between adjacent open cells). Use BFS to find a shortest path from S→E, and output the path either as a list of coordinates or directional steps.
- 3 (*Optional*) Compare DFS vs BFS on the same maze: implement a DFS-based solver and observe whether it always finds the shortest path.

Outcome: Bridge puzzles to graph algorithms, reinforcing that both recursion (DFS) and queue-based search (BFS) generalize across problem domains.

Group Challenge — Extensions

After three days of “learning by doing” (N-Queens, Sudoku, Word Ladders, Mazes), today we’ll survey several classic puzzles and their high-level algorithms. You won’t write full code now, but you should:

1. See how each problem is posed.
2. Follow the pseudocode sketch.
3. Think of one small way you might extend or improve it.

Puzzles & Pseudocode Stubs

- **Lights Out**

A grid of lights toggles itself and its orthogonal neighbors when you press a cell. Goal: turn all lights off.

```
function solveLightsOut(state):
    if all cells are OFF:
        return solution # base case
    pick a cell (r,c)
    toggle(state, r, c) # flip this cell + neighbors
    if solveLightsOut(state):
        return solution
    toggle(state, r, c) # backtrack
    return failure
```

- **Josephus**

n people stand in a circle; every kth person is eliminated until one remains.

```
function josephus(n, k):
    if n == 1:
        return 0 # survivor index in zero-based
    prev = josephus(n-1, k)
    return (prev + k) mod n # re-index into current circle
```

- **Tower of Hanoi**

Move n disks among three pegs, one disk at a time, never placing larger atop smaller.

```
function hanoi(n, src, aux, dest):
    if n == 1:
        move disk from src to dest
        return
    hanoi(n-1, src, dest, aux) # move top n-1 to auxiliary
    move disk from src to dest
    hanoi(n-1, aux, src, dest) # move n-1 from auxiliary to dest
```

- **Coin Change (Dynamic)**

Count ways to make amount X using unlimited coins of given denominations.


```

function countWays(coins, X):
    dp = array[0..X] with dp[0] = 1
    for each coin in coins:
        for amt = coin to X:
            dp[amt] += dp[amt - coin]
    return dp[X]

```

- **Word Search**

Given a grid of letters, determine if a target word exists by moving through adjacent cells (up/down/left/right) without reuse.

```

function exists(word, grid):
    for each cell (r,c):
        if dfs((r,c), 0):
            return TRUE
    return FALSE

function dfs(pos, i):
    if i == length(word):
        return TRUE # full word found
    if grid[pos] != word[i]:
        return FALSE
    mark pos visited
    for each neighbor n of pos:
        if not visited(n) and dfs(n, i+1):
            return TRUE
    unmark pos
    return FALSE

```

Exercises

1. **Trace Lights Out.** On this 3×3 sample, show the first two presses (cells) you'd choose to begin solving:

```

1 0 0
0 1 0
0 0 1

```

2. **Compute Josephus.** For $n=7$, $k=3$, who survives? Show your calls and return values from the recurrence.
3. **(Optional)** Pick one of today's puzzles and sketch a small extension. For instance:

- Lights Out on a torus
- Variable k in Josephus
- Tower of Hanoi variants
- Greedy heuristics for Coin Change
- Word Search backtracking

Be ready to describe your idea in two sentences.

Outcome: Familiarize everyone with a broader family of recursive, dynamic-programming, and graph-search problems—and give you a chance to look under the hood of each algorithm, even if you’re not writing full code today.

Queues – Real-World Simulation

Problem Description:

Model a bank teller’s line as a queue: customers arrive, wait in FIFO order, and are served one at a time. We’ll also demo a stack-based undo/redo of teller actions.

Live-Code: Bank-Line Simulator

Use a queue to enqueue arriving customers and dequeue for service.

Stack Demo: Undo/Redo

Record each served customer on an `undoStack`. Allow “undo” (move from `undoStack` → `redoStack`) and “redo” (move back).

Pseudocode Stub

```
function serveBank(arrivals):
    Q = empty queue
    for customer in arrivals:
        enqueue(Q, customer)

    while not empty(Q):
        c = dequeue(Q)
        serve(c)
        push(undoStack, c)

function undo():
    if not empty(undoStack):
        c = pop(undoStack)
        unserve(c)
        push(redoStack, c)

function redo():
    if not empty(redoStack):
        c = pop(redoStack)
        serve(c)
        push(undoStack, c)
```

Visual Walkthrough

A quick state-diagram table (front→back):

Step	Queue (front→back)	undoStack	redoStack
start	[A, B, C, D]	[]	[]
serve A	[B, C, D]	[A]	[]
serve B	[C, D]	[A, B]	[]
undo	[C, D]	[A]	[B]
redo	[C, D]	[A, B]	[]

Complexity Note

- All queue operations (‘enqueue’/‘dequeue’) and stack operations (‘push’/‘pop’) run in $O(1)$ time per action.

Sample Input / Scenario

- Arrival order: [Alice, Bob, Carol, Dave]
- Actions:
 - Serve Alice, Bob, Carol
 - Undo twice
 - Redo once
- Track queue contents and stack states after each step.

Exercises

1. Given arrivals [Alice, Bob, Carol], serve two, undo one, then list:
 - Remaining queue
 - Contents of `undoStack` and `redoStack`
2. Extend the stub to support “VIP” arrivals that jump to the front of the queue. Sketch your code change.

Outcome:

Apply queues and stacks in a realistic simulation—reinforcing FIFO service and undo/redo mechanics.

Game AI – Tic-Tac-Toe

Problem Description:

Build an AI opponent for Tic-Tac-Toe using a decision-tree search (*minimax*). The AI enumerates possible move sequences and chooses the optimal one.

Pseudocode Stub

```
function minimax(board, player):
    if terminal(board):
        return utility(board)          # win / loss / draw score

    if player == MAX:                  # MAX = 'O'
        bestVal = -infty
        for move in legalMoves(board):
            val = minimax(apply(board, move), MIN)
            bestVal = max(bestVal, val)
        return bestVal
    else:                              # MIN = 'X'
        bestVal = +infty
        for move in legalMoves(board):
            val = minimax(apply(board, move), MAX)
            bestVal = min(bestVal, val)
        return bestVal

function bestMove(board, player):
    best = null
    bestVal = (player==MAX ? -infty : +infty)
    for move in legalMoves(board):
        val = minimax(apply(board, move), opposite(player))
        if (player==MAX and val > bestVal) or
            (player==MIN and val < bestVal):
            bestVal = val
            best = move
    return best
```

Exercises

1. **Trace minimax.** On the board below, X (= MIN) has just played; O (= MAX) has *one* legal move left—the centre square. Show the call sequence (*minimax* → *utility*) and the utility value finally returned.

X	O	X
O	–	O
X	O	X

2. **Implement utility.** Sketch pseudocode for `utility(board)` so that it returns `+1` if O wins, `–1` if X wins, and `0` for a draw.

Outcome:

Expose the mechanics of recursive decision-tree search and utility evaluation in simple game AI.

Event-Driven Modeling: Traffic / Queue

Problem Description:

Model a traffic-light controller (or bank queue) as an *event-driven* system. Cars/lights fire *events*—ARRIVAL or DEPARTURE—and a central loop dispatches them in time order.

Live Code – Event-Loop Simulator

Demonstrate an **event priority queue**¹ that holds pending events while a dispatcher invokes handlers.

Pseudocode Stub

```
function runSimulation(initialEvents):
    Q = empty priority-queue          # ordered by (time, D<A)
                                     # no built-in priority flag yet
    waitingList = []                 # current cars in line
    queueLength = 0                  # global length variable
    nextID = 1 + max(id for (_,_,id) in initialEvents)

    for e in initialEvents:
        enqueue(Q, e)

    while not empty(Q):
        (t, type, cid) = dequeue(Q)  # earliest event
        handleEvent(type, t, cid)

function handleEvent(type, t, cid):
    global nextID, queueLength

    if type == ARRIVAL:
        waitingList.append(cid)
        queueLength = len(waitingList)

        if queueLength == 1:          # line was empty
            schedule(DEPARTURE, t + serviceTime(), cid)

        schedule(ARRIVAL, t + arrivalInterval(), nextID)
        nextID += 1

    elif type == DEPARTURE:
        waitingList.pop(0)             # front car leaves
        queueLength = len(waitingList)

        if queueLength > 0:           # next car departs later
            nextFront = waitingList[0]
            schedule(DEPARTURE, t + serviceTime(), nextFront)
```

¹Tie rule: Events are ordered by (time, type) with DEPARTURE < ARRIVAL, i.e., departure events come before arrival events in the priority queue and are thus considered higher-priority.

```

updateSystemState(t, type, cid, queueLength)

function schedule(type, tNew, cid):
    event = (tNew, type, cid)    # plain 3-field tuple, ordered by time and type
    enqueue(Q, event)           # O(log n) priority-queue insert

function arrivalInterval(): return ⟨user-supplied⟩    # e.g. 3 s
function serviceTime():    return ⟨user-supplied⟩    # e.g. 4 s

function updateSystemState(t, type, cid, queueLen):
    log(t, type, cid, queueLen)    # placeholder|stats / GUI / etc.

```

Sample Input / Scenario

- **Warm-up** events: $[(0, \text{ARRIVAL}, c_1), (0, \text{ARRIVAL}, c_2)]$
- **Exercise set** arrivals: $[0, 2, 5]$
- **Handlers**
 - **ARRIVAL** – enqueue next arrival, *increment* queueLength
 - **DEPARTURE** – enqueue next departure, *decrement* queueLength
- Track both the *event queue* Q and the *system state* ($\text{queueLength} = |\text{waitingList}|$) at every step.

Exercises

1. **Trace event loop.** Assume $\text{arrivalInterval} = 6\text{ s}$ and $\text{serviceTime} = 4\text{ s}$. For arrivals at $[0, 2, 5]$, list the first five rows of

$$(\text{step}, Q, \text{dequeued } (t, \text{event}), \text{waitingList}, \text{queueLength})$$
 and write a short note for each row.
2. **PRIORITY event.** Extend `schedule` to accept a priority flag (e.g. extra tuple field) so that a “PRIORITY” event jumps to the front.

Outcome:

Understand event-driven architecture—event queues, dispatch loops, dynamic scheduling, and real-time queue-length tracking.

Ideation Workshop

Problem Description

Working in small teams, you will *discover* and *define* a real bottleneck on campus or in the local community—e.g. library room-booking clashes, canteen queue congestion, or hostel laundry scheduling. The aim is to express the problem precisely and decide which issues merit an algorithmic solution.

Live Demonstration – Minimal Vote-Tally Script

The facilitator walks through a 10-line Python function that counts dot-votes and produces a ranked list of ideas.

```
# ideas    : list[str]
# ballots  : list[list[int]] # indices voted by each student
# Each student can cast up to 3 dots.

def tally_votes(ideas, ballots, dots_per_student=3):
    scores = {i: 0 for i in range(len(ideas))}
    for ballot in ballots:
        assert len(ballot) <= dots_per_student
        for idx in ballot:
            scores[idx] += 1      # one dot = one point
    return sorted(scores.items(), key=lambda p: -p[1])
```

Pseudocode Stub – Workshop Helper

```
function runIdeationWorkshop(studentIdeas):
    clusters = cluster_similar(studentIdeas)    # manual grouping / DBSCAN
    display_table(clusters)                    # show on projector

    votes    = collect_dot_votes(clusters)      # 3 dots per learner
    ranking  = tally_votes(clusters, votes)     # reuse the live-deno helper

    topIdea  = ranking[0]
    return draft_problem_statement(topIdea)
```

Warm-Up – Everyday Bottlenecks

Discuss in pairs: which campus processes could benefit from an algorithmic or data-driven approach?

1. Library *seat* / *room* booking clashes
2. Canteen queue congestion during peak breaks
3. Hostel laundry-machine scheduling fairness
4. Event ticket allocation transparency

Guided Brainstorm (15 min)

- **Individually:** Jot down as many campus/community problems as you can.
- **Teams of 3–4:** Cluster similar ideas and assign a concise label.

Structured Voting (5 min)

Each participant receives **three** dots/stickers. Distribute them among the clusters you believe are most pressing.

Idea / Cluster	Votes

Team Deliverable (10 min)

For the top-ranked idea, craft a concise problem statement covering:

- **Problem** – What exactly happens and why is it undesirable?
- **Stakeholders** – Who is affected and how?
- **Evidence** – A metric or anecdote illustrating scale or urgency.
- **Impact Goal** – Define success in measurable terms.

Outcome:

Broaden perspective beyond individual algorithms, nurturing interdisciplinary vision and opportunity discovery.

Mini-Pitch – Reflection

Activity Overview

In teams of 3–4, present a concise, 3-minute pitch of your Day 9 solution (“ID-Verification Chaos” or other top idea). Focus on clarity and impact.

Pitch Guidelines

- **Problem Summary:** State the core issue in one sentence.
- **Stakeholders & Evidence:** Highlight who is affected and key data.
- **Solution Sketch:** Describe your variables or algorithmic approach.
- **Impact Metrics:** Specify how success will be measured.
- **Time Limit:** 3 minutes exactly (use a visible timer).

Reflection Prompts

After all pitches, take 10 minutes for individual written reflection:

1. What was the most challenging aspect of defining the problem?
2. Which component (stakeholder insight, evidence, or solution sketch) do you value most, and why?
3. How would you refine your approach or pitch next time?
4. What new perspectives did you gain from listening to peers?

Team Deliverable

Submit the following at the end of the session:

- A one-page bullet summary of your pitch (max 200 words).
- Written answers to the four reflection prompts.

Outcome:

Build confidence in presenting solutions and prepare for Phase 2 capstone challenges.